

Semi-Dynamic Variability

Between Static Configuration and Dynamic Reconfiguration

Christian Dietrich

13. September 2021



Special-Purpose
Systems



Different
Platforms



Security-
Updates



Dynamic
Scalability



Debugging &
Introspection



Extendable and/or
Reducible
Feature Set





Special-Purpose
Systems



Different
Platforms



Security-
Updates



Dynamic
Scalability



Debugging &
Introspection



Extendable and/or
Reducible
Feature Set





Special-Purpose
Systems

Different
Platforms

Security-
Updates

Dynamic
Scalability

Debugging &
Introspection



Extendable and/or
Reducible
Feature Set



Special-Purpose
Systems

Different
Platforms

Security-
Updates

Dynamic
Scalability

Debugging &
Introspection



Extendable and/or
Reducible
Feature Set



Special-Purpose
Systems

Different
Platforms

Security-
Updates

Dynamic
Scalability

Debugging &
Introspection

Extendable and/or
Reducible
Feature Set





Variable System Software is Adaptable System Software



Software Product Line





Variable System Software is Adaptable System Software



Software Product Line





Static and Dynamic Variability

Creation: of a variant's binary form

Binding: Integration of Variant into Program



Creation: of a variant's binary form

Binding: Integration of Variant into Program

```
#ifdef CONFIG_SMP
irq_disable();
spin_acquire(&lock);
#else
irq_disable();
#endif
```

Static Variability

- Creating: Compile-Time
- Binding: Compile-Time
- + No Run-Time Overhead
- Each combination: A different binary



Creation: of a variant's binary form

Binding: Integration of Variant into Program

```
#ifdef CONFIG_SMP
irq_disable();
spin_acquire(&lock);
#else
irq_disable();
#endif
```

```
if (config_smp) {
    irq_disable();
    spin_acquire(&lock);
} else {
    irq_disable();
}
```

Static Variability

- Creating: Compile-Time
- Binding: Compile-Time
- + No Run-Time Overhead
- Each combination: A different binary

Dynamic Variability

- Creation: Hybrid Variants
- Binding: Time-of-Use
- Repeated binding costs
- + Flexible and reconfigurable



Creation: of a variant's binary form

Binding: Integration of Variant into Program

```
#ifdef CONFIG_SMP
irq_disable();
spin_acquire(&lock);
#else
irq_disable();
#endif
```

```
if (config_smp) {
    irq_disable();
    spin_acquire(&lock);
} else {
    irq_disable();
}
```

Static Variability

- Creating: Compile-Time
- Binding: Compile-Time
- + No Run-Time Overhead
- Each combination: A different binary

Dynamic Variability

- Creation: Hybrid Variants
- Binding: Time-of-Use
- Repeated binding costs
- + Flexible and reconfigurable

~> Semi-Dynamic Variability

Static Creation and **efficient** dynamic (re-)binding



Multiverse

Efficient Semi-Dynamic Variability

EuroSys'19

- Preparing variants at compile time
- Efficient binding by run-time code patching



Wait-free code patching

Run-time code patching without interruption

OSDI'20

- Preparing changes in a separate address space
- Incremental migration of threads



Multiverse

Efficient Semi-Dynamic Variability

EuroSys'19

- Preparing variants at compile time
- Efficient binding by run-time code patching



Wait-free code patching

Run-time code patching without interruption

OSDI'20

- Preparing changes in a separate address space
- Incremental migration of threads



Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```



Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

GCC

Code Segment

```
spin_irq_lock:  
    cmp    <smp>, 0  
    je     .else  
    cli  
    call   spin_acquire  
    ret  
.else:  
    cli  
    ret
```




Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

+ Multiverse
GCC

Code Segment

```
spin_irq_lock.smp=1:  
cli  
call spin_acquire  
ret
```

```
spin_irq_lock.smp=0:  
cli  
ret
```

```
spin_irq_lock:  
cmp    <smp>, 0  
je     .else  
cli  
call   spin_acquire  
ret  
.else:  
cli  
ret
```



Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

var

func

callsite

+ Multiverse
GCC

Multiverse
Deskriptoren

Code Segment

```
spin_irq_lock.smp=1:  
    cli  
    call spin_acquire  
    ret
```

```
spin_irq_lock.smp=0:  
    cli  
    ret
```

```
spin_irq_lock:  
    cmp    <smp>, 0  
    je     .else  
    cli  
    call  spin_acquire  
    ret  
.else:  
    cli  
    ret
```



Initial geladenes Code Segment

```
foo:
    ...
    call    multiverse_commit
    ...
    call    spin_irq_lock
    ...
    ret
```

```
spin_irq_lock.smp=1:
    cli
    call    spin_acquire
    ret
```

```
spin_irq_lock.smp=0:
    cli
    ret
```

```
spin_irq_lock:
    cmp     <smp>, 0
    je      .else
    cli
    call    spin_acquire
    ret
.else:
    cli
    ret
```

Multiverse
Deskriptoren



Initial geladenes Code Segment

```
foo:
    ...
    call multiverse_commit
    ...
    call spin_irq_lock
    ...
    ret
```

```
spin_irq_lock.smp=1:
    cli
    call spin_acquire
    ret
```

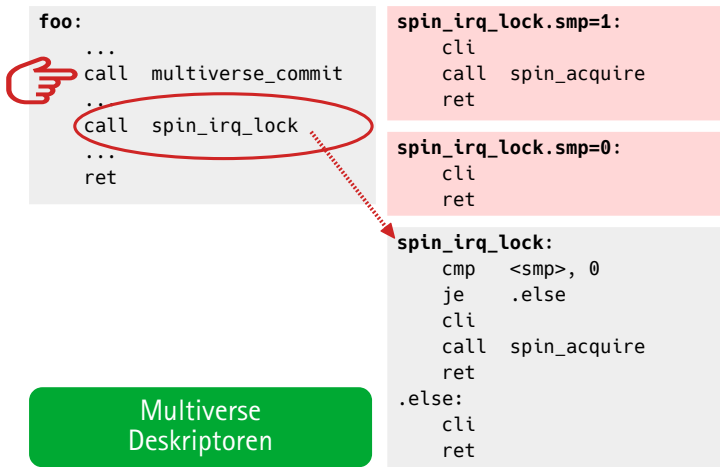
```
spin_irq_lock.smp=0:
    cli
    ret
```

```
spin_irq_lock:
    cmp    <smp>, 0
    je     .else
    cli
    call   spin_acquire
    ret
.else:
    cli
    ret
```

Multiverse
Deskriptoren



Initial geladenes Code Segment





Patched (smp == 1)

```
foo:
...
call multiverse_commit
...
call spin_irq_lock.smp=1
...
ret
```

```
spin_irq_lock.smp=1:
cli
call spin_acquire
ret
```

```
spin_irq_lock.smp=0:
cli
ret
```

```
spin_irq_lock:
jmp spin_irq_lock.smp=1
je .else
cli
call spin_acquire
ret
.else:
cli
ret
```

Multiverse
Deskriptoren



Patched (smp == 0)

```
foo:
  ...
  call multiverse_commit
  ...
  cli
  ...
  ret
```

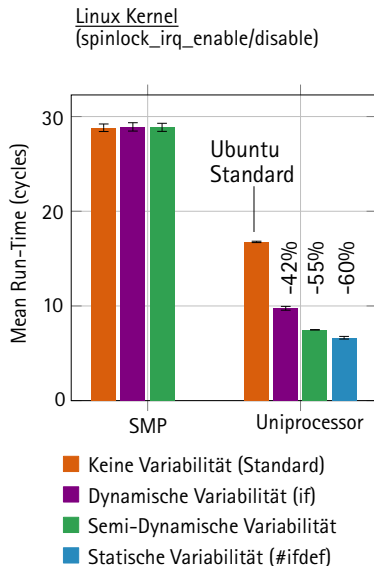
Call-Site Inlining!

```
spin_irq_lock.smp=1:
  cli
  call spin_acquire
  ret
```

```
spin_irq_lock.smp=0:
  cli
  ret
```

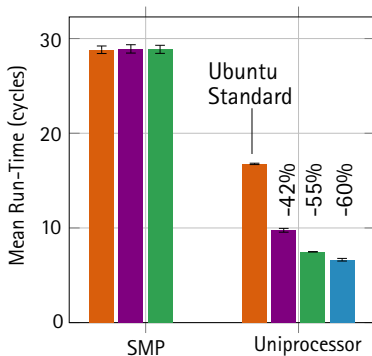
```
spin_irq_lock:
  jmp spin_irq_lock.smp=0
  je .else
  cli
  call spin_acquire
  ret
.else:
  cli
  ret
```

Multiverse
Deskriptoren



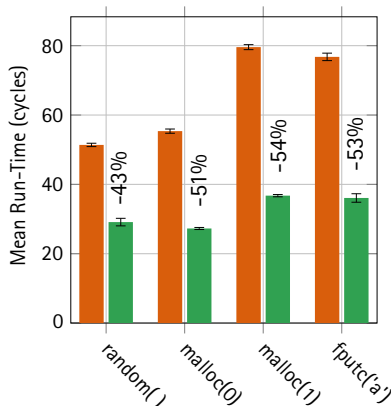


Linux Kernel
(spinlock_irq_enable/disable)



- Keine Variabilität (Standard)
- Dynamische Variabilität (if)
- Semi-Dynamische Variabilität
- Statische Variabilität (#ifdef)

Musl C Bibliothek
(Single Threaded Modus)

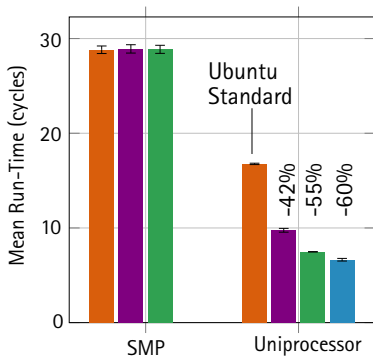


- Keine Variabilität (Standard)
- Semi-Dynamische Variabilität



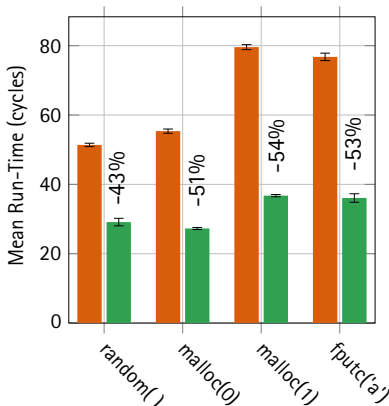
Lock Elision in Kernel 4.16:

- 1161 spin-lock call sites
- +40 KiB size (zipped total: 10 MiB)



- Keine Variabilität (Standard)
- Dynamische Variabilität (if)
- Semi-Dynamische Variabilität
- Statische Variabilität (#ifdef)

Musl C Bibliothek (Single Threaded Modus)



- Keine Variabilität (Standard)
- Semi-Dynamische Variabilität



Lock Elision in Kernel 4.16:

- 1161 spin-lock call sites
- +40 KiB size (zipped total: 10 MiB)

Musl C Bibliothek
(Single Threaded Modus)

Multiverse – Efficient Semi-Dynamic Variability

- + Variability specification within language
- + Run-time cost like static variability
- + Flexibility like dynamic variability

0 SMP Uniprocessor

- Keine Variabilität (Standard)
- Dynamische Variabilität (if)
- Semi-Dynamische Variabilität
- Statische Variabilität (#ifdef)

0 random() malloc(0) malloc(1) fputc('a')

- Keine Variabilität (Standard)
- Semi-Dynamische Variabilität

-53%



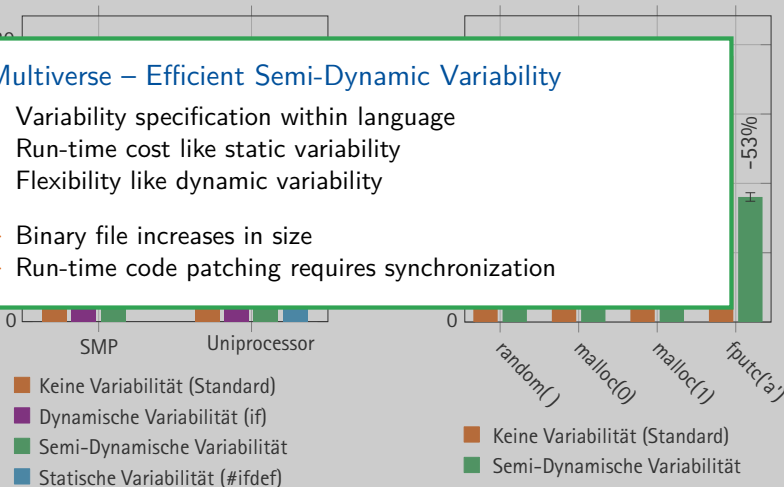
Lock Elision in Kernel 4.16:

- 1161 spin-lock call sites
- +40 KiB size (zipped total: 10 MiB)

Musl C Bibliothek
(Single Threaded Modus)

Multiverse – Efficient Semi-Dynamic Variability

- + Variability specification within language
- + Run-time cost like static variability
- + Flexibility like dynamic variability
- Binary file increases in size
- Run-time code patching requires synchronization





■ Example Source Code

- `git clone https://github.com/luhsra/multiverse-atlas-demo.git`
- `./init =>` Clone dependent Libraries
- You need a mmview-enabled kernel for example 02/03

■ Virtual Machine Image

- Download:
`https://lab.sra.uni-hannover.de/Publications/2021/netsys-tutorial/`
- Start: `qemu-system-x86_64 -m 4G -enable-kvm -hda hda.qcow2 -smp 4`
- User/Passwort: root/root, user/user

■ Running Virtual Machine (temporary)

- `ssh user@lab.sra.uni-hannover.de -p 2250`
- Password: FIXME
- Machine will be available only for the time of the tutorial.



Goal

Demonstrate Semi-Dynamic Lock-Elision with Multiverse

- Static Creation with `__attribute__((multiverse))`
 - Multiverse Variables (MV) are „features“. Default Range: {0, 1}
 - Attributed functions are multiversed at MV usages.
 - Cross-Product of MV ranges.
- Binary File: 01-multiverse
 - Regular functions: `lock()`, `unlock()`
 - Use `objdump -d 01-multiverse` to inspect assembler
 - Search for `lock.multiverse.smp_0`

■ Multiverse Metadata

```
fn: lock 0x55bd78dc92e0, 2 variants
mvfn: 0x55bd78dc9410 (vars 1)
  assign: smp in [0, 0]
mvfn: 0x55bd78dc9420 (vars 1)
  assign: smp in [1, 1]
```



```
Without Multiverse (smp=0)
work() -> 83.376870 ms
work() -> 70.550493 ms
work() -> 70.670475 ms
work() -> 70.578593 ms
work() -> 73.273835 ms
Average (n=5): 73.690053
```

```
Without Multiverse (smp=1)
work() -> 511.101093 ms
work() -> 511.533894 ms
work() -> 512.103367 ms
work() -> 519.490820 ms
work() -> 511.268393 ms
Average (n=5): 513.099513
```

```
With Multiverse (smp=0)
work() -> 18.856244 ms
work() -> 18.820774 ms
work() -> 18.878794 ms
work() -> 18.757183 ms
work() -> 18.795324 ms
Average (n=5): 18.821664
```

```
With Multiverse (smp=1)
work() -> 509.799176 ms
work() -> 509.639876 ms
work() -> 518.221404 ms
work() -> 511.009342 ms
work() -> 510.144087 ms
Average (n=5): 511.762777
```



Multiverse

Efficient Semi-Dynamic Variability

EuroSys'19

- Preparing variants at compile time
- Efficient binding by run-time code patching



Wait-free code patching

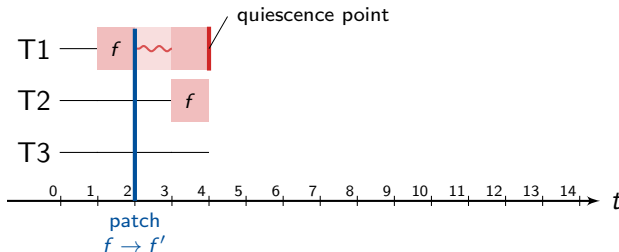
Run-time code patching without interruption

OSDI'20

- Preparing changes in a separate address space
- Incremental migration of threads

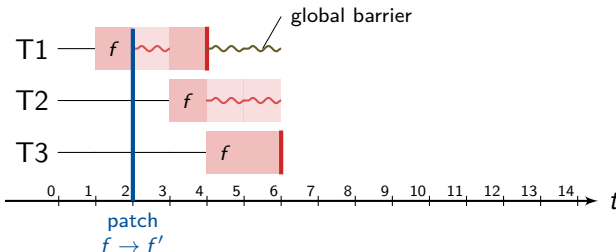


Run-Time Code Patching requires Quiescence!



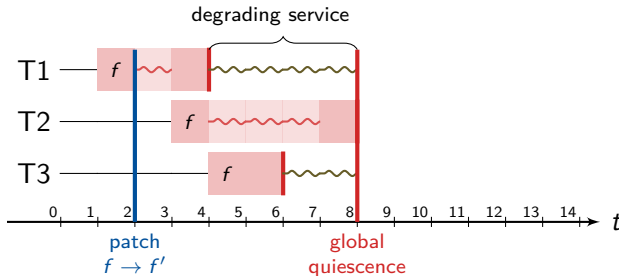


Run-Time Code Patching requires Quiescence!



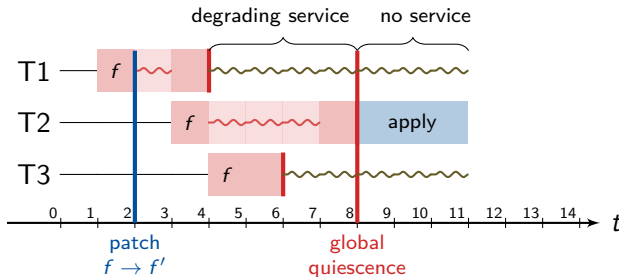


Run-Time Code Patching requires Quiescence!



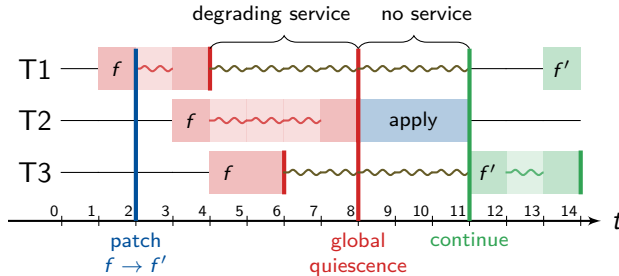
- Global barrier bears risk of dead-lock
- Impact on service quality

Run-Time Code Patching requires Quiescence!



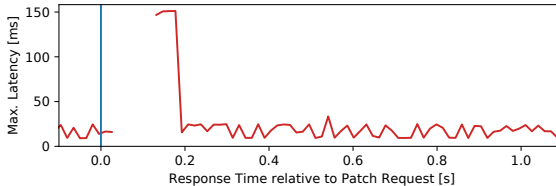
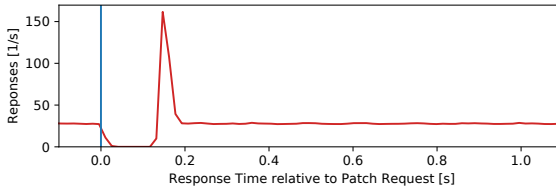
- Global barrier bears risk of dead-lock
- Impact on service quality
- + Strong consistency through global quiescence

Run-Time Code Patching requires Quiescence!



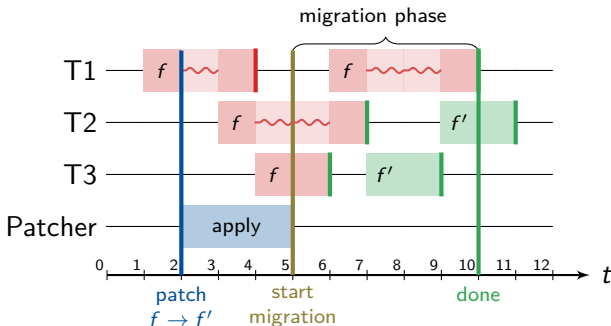
- Global barrier bears risk of dead-lock
- Impact on service quality
- + Strong consistency through global quiescence

OpenLDAP



- Glo
- Imp

+ Strong consistency through global quiescence



Wait-Free Code Patching

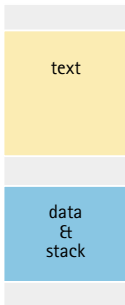
Incremental Application of Patch

- Prepare the modified code segment
- Threads migrate themselves and without synchronization
- Old and new variants co-exist (for shorter or longer periods of time)



address
space

Generation 0

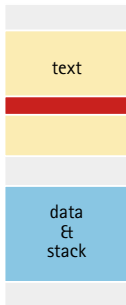




Address-Space Views

address
space

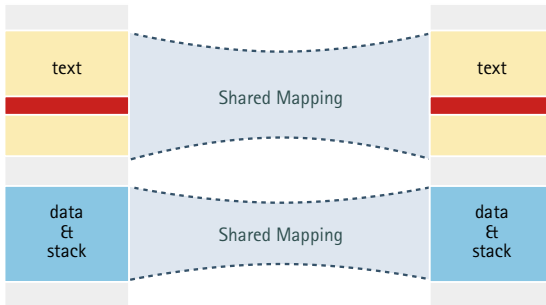
Generation 0

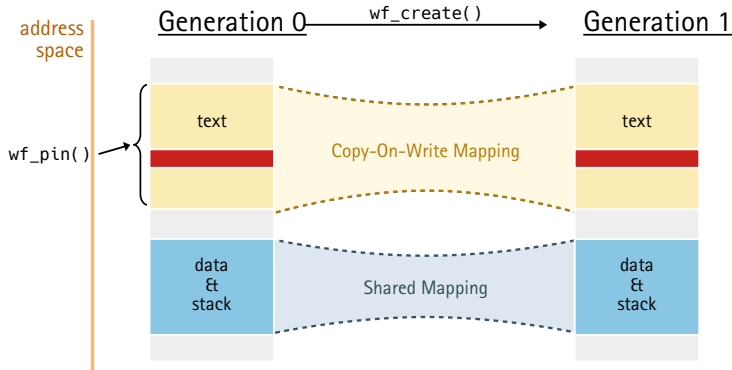


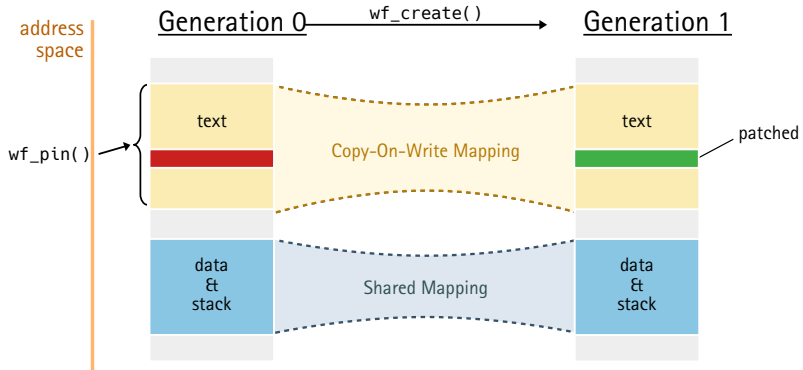


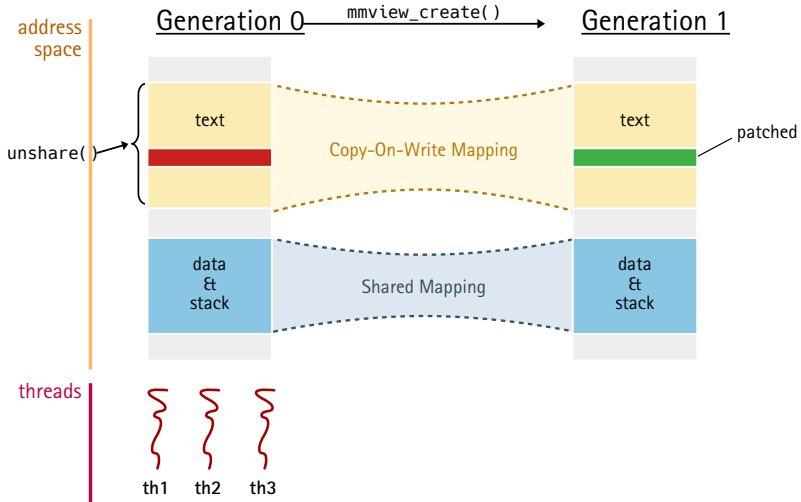
address
space

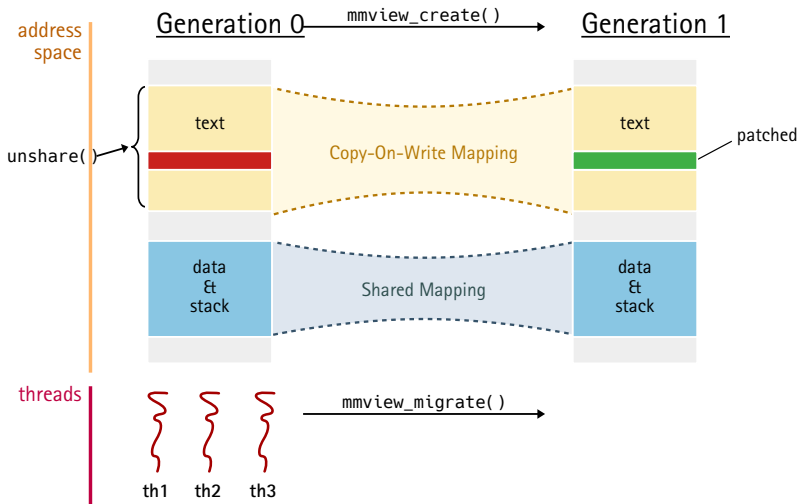
Generation 0 $\xrightarrow{\text{wf_create()}}$ Generation 1

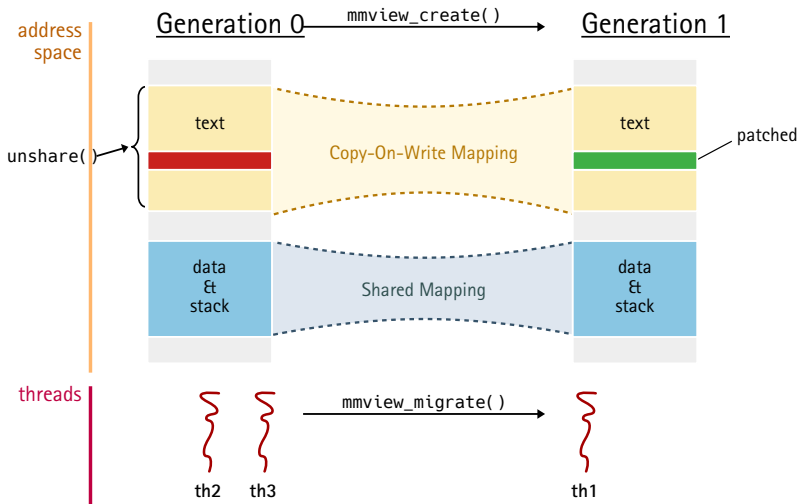


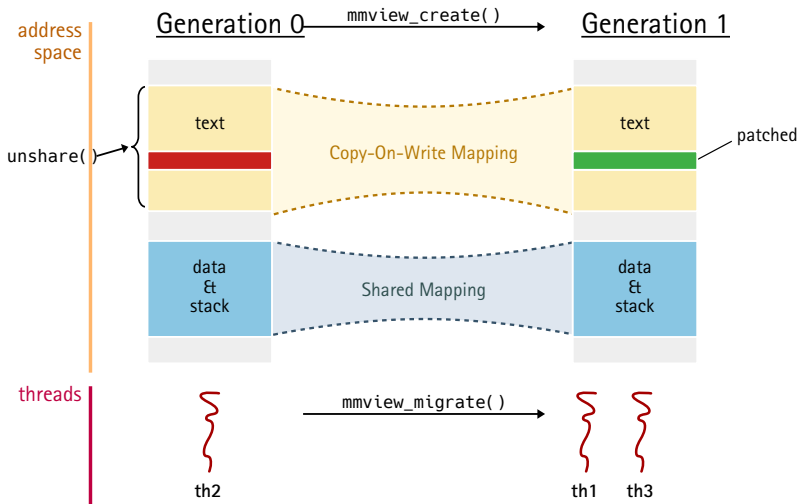


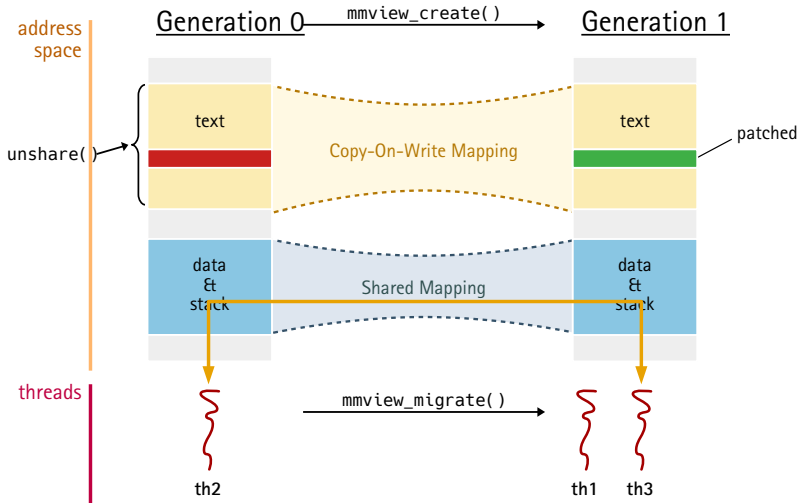






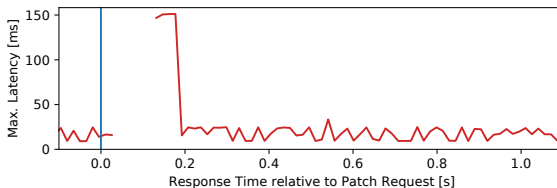
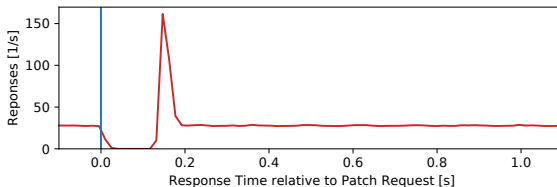








OpenLDAP



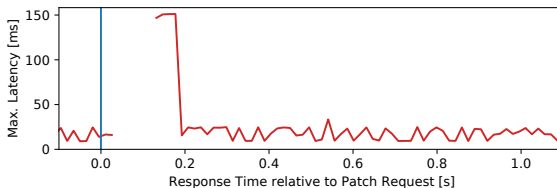
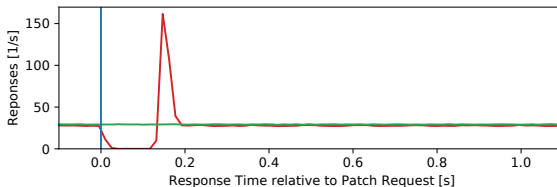
th2

th1

th3



OpenLDAP



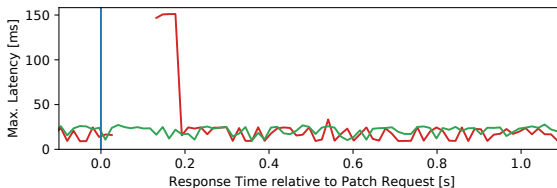
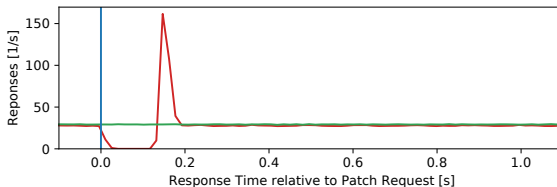
th2

th1

th3



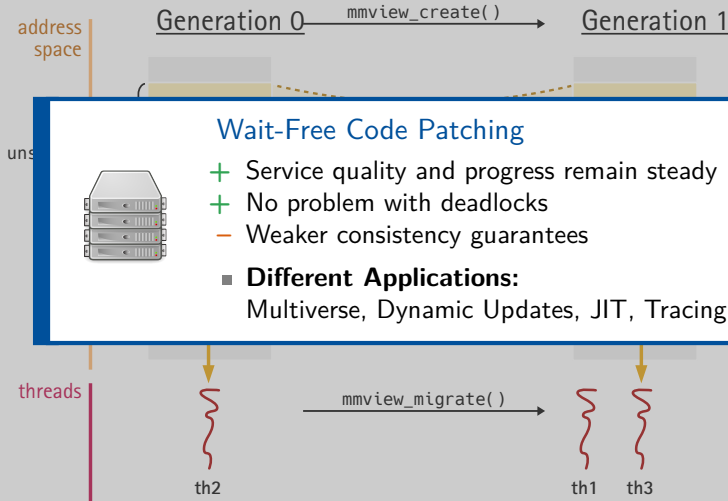
OpenLDAP



th2

th1

th3





Goal

Demonstrate Address-Space View Migration with Multiverse

```
mmview_t view = mmview_create();  
mmview_migrate(view);  
  
multiversed_var = 1;  
multiverse_commit();  
desired_view = view;
```

- Create new address-space view and migrate the current thread.
- Commit current MV state into text segment.
- Warning: multiverse_var is shared!

```
if (mmview_current() != desired_view) {  
    mmview_migrate(desired_mmview);  
}
```

Each thread checks at its local-quiescence point, whether it is already in the correct view. If not
→ `mmview_migrate()`;



```
thread2 -> view B
thread0 -> view B
mmview_create() -> 1
mmview_migrate(1) -> 0
thread1 -> view B
thread1 migrate: mmview_migrate(1) -> 0
thread2 -> view B
thread2 migrate: mmview_migrate(1) -> 0
thread1 -> view A
thread0 -> view B
thread0 migrate: mmview_migrate(1) -> 0
thread1 -> view A
thread0 -> view A
thread2 -> view A
thread1 -> view A
thread0 -> view A
thread1 -> view A
thread1 -> view A
```



Goal

Demonstrate concurrent existence of multiversed address-space views.

- **Problem:** Profiling in multi-threaded processes is slow!
 - One (shared!) counter per profiling point.
 - Massive cache line bouncing between threads.
 - Your Choice: Racy counter updates OR run-time impact of atomic operations
- **Approach:** Let only one thread profile at once
 - Two Views with Multiverse: Profiling and Non-Profiling View
 - Let N threads (probabilistically) switch to the profiling view
 - Control parallelism in profiling with semaphore



■ 4 threads calculating, no profiling

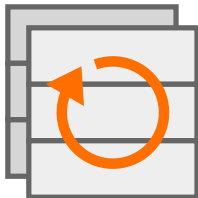
```
./03-concurrent 4 0  
Avg. time for fib(30) = 0.956386 ms; -nan branches/call 400
```

■ 4 threads calculating, 4 threads profiling

```
./03-concurrent 4 4  
Avg. time for fib(30) = 106.137128 ms; 0.509816 branches/call 400
```

■ 4 threads calculating, 1 threads profiling

```
./03-concurrent 4 1  
Avg. time for fib(30) = 1.088290 ms; 0.500000 branches/call 400
```



*Operating
System
Group*

- Operating System Group: Young Basic Research Group
 - Topics: Non-Volatile Memory, Variability, Dependable Real-Time Systems
 - Techniques: Virtual memory, Compilers, Fault Injection
 - Open PhD Position: 100% E13, 3 years
- Technical University Hamburg
 - ~8000 students, ~90 professors
 - Strong systems-oriented computer science
 - South of the Elbe

⇒ Get in contact, if interested



Multiverse

Efficient Semi-Dynamic Variability

EuroSys'19

- Preparing variants at compile time
- Efficient binding by run-time code patching



Wait-free code patching

Run-time code patching without interruption

OSDI'20

- Preparing changes in a separate address space
- Incremental migration of threads



- [1] Florian Rommel, Christian Dietrich, Daniel Friesel u. a. „From Global to Local Quiescence: Wait-Free Code Patching of Multi-Threaded Processes“. In: *14th Symposium on Operating System Design and Implementation (OSDI '20)*. Nov. 2020, S. 651–666. URL: <https://www.usenix.org/conference/osdi20/presentation/rommel>.
- [2] Florian Rommel, Christian Dietrich, Michael Rodin u. a. „Multiverse: Compiler-Assisted Management of Dynamic Variability in Low-Level System Software“. In: *Fourteenth EuroSys Conference 2019 (EuroSys '19)* (Dresden, Germany). New York, NY, USA: ACM Press, 2019. ISBN: 978-1-4503-6281-8. DOI: 10.1145/3302424.3303959.
- [3] Florian Rommel, Lennart Glauer, Christian Dietrich u. a. „Wait-Free Code Patching of Multi-Threaded Processes“. In: *Proceedings of the 10th SOSP Workshop on Programming Languages and Operating Systems (PLOS '19)*. Huntsville, ON, Canada: ACM, 2019, S. 23–29. ISBN: 978-1-4503-7017-2. DOI: 10.1145/3365137.3365404.
- [4] Valentin Rothberg, Christian Dietrich, Alexander Graf u. a. „Function Multiverses for Dynamic Variability“. In: *Foundations and Applications of Self* Systems*. Hrsg. von Jesper Andersson, Rafael Capilla und Holger Eichelberger. Augsburg, 2016. URL: https://www4.cs.fau.de/Publications/2016/rothberg_16_dspl.pdf.