

# Semi-Dynamic Variability

Between Static Configuration and Dynamic Reconfiguration

Christian Dietrich

13. September 2021



Special-Purpose  
Systems



Different  
Platforms



Security-  
Updates



Dynamic  
Scalability



Debugging &  
Introspection



Extendable and/or  
Reducible  
Feature Set





Special-Purpose  
Systems



Different  
Platforms



Security-  
Updates



Dynamic  
Scalability



Debugging &  
Introspection



Extendable and/or  
Reducible  
Feature Set





Special-Purpose  
Systems

Different  
Platforms

Security-  
Updates

Dynamic  
Scalability

Debugging &  
Introspection



Extendable and/or  
Reducible  
Feature Set



Special-Purpose  
Systems

Different  
Platforms

Security-  
Updates

Dynamic  
Scalability

Debugging &  
Introspection



Extendable and/or  
Reducible  
Feature Set



Special-Purpose  
Systems

Different  
Platforms

Security-  
Updates

Dynamic  
Scalability

Debugging &  
Introspection

Extendable and/or  
Reducible  
Feature Set





# Variable System Software is Adaptable System Software



## Software Product Line





# Variable System Software is Adaptable System Software



## Software Product Line







# Static and Dynamic Variability

**Creation:** of a variant's binary form

**Binding:** Integration of Variant into Program



**Creation:** of a variant's binary form

**Binding:** Integration of Variant into Program

```
#ifdef CONFIG_SMP
irq_disable();
spin_acquire(&lock);
#else
irq_disable();
#endif
```

## Static Variability

- Creating: Compile-Time
- Binding: Compile-Time
- + No Run-Time Overhead
- Each combination: A different binary



**Creation:** of a variant's binary form

**Binding:** Integration of Variant into Program

```
#ifdef CONFIG_SMP
irq_disable();
spin_acquire(&lock);
#else
irq_disable();
#endif
```

```
if (config_smp) {
    irq_disable();
    spin_acquire(&lock);
} else {
    irq_disable();
}
```

## Static Variability

- Creating: Compile-Time
- Binding: Compile-Time
- + No Run-Time Overhead
- Each combination: A different binary

## Dynamic Variability

- Creation: Hybrid Variants
- Binding: Time-of-Use
- Repeated binding costs
- + Flexible and reconfigurable



**Creation:** of a variant's binary form

**Binding:** Integration of Variant into Program

```
#ifdef CONFIG_SMP
irq_disable();
spin_acquire(&lock);
#else
irq_disable();
#endif
```

```
if (config_smp) {
    irq_disable();
    spin_acquire(&lock);
} else {
    irq_disable();
}
```

## Static Variability

- Creating: Compile-Time
- Binding: Compile-Time
- + No Run-Time Overhead
- Each combination: A different binary

## Dynamic Variability

- Creation: Hybrid Variants
- Binding: Time-of-Use
- Repeated binding costs
- + Flexible and reconfigurable

## ~> Semi-Dynamic Variability

Static Creation and **efficient** dynamic (re-)binding



## Multiverse

Efficient Semi-Dynamic Variability

*EuroSys'19*

- Preparing variants at compile time
- Efficient binding by run-time code patching



## Wait-free code patching

Run-time code patching without interruption

*OSDI'20*

- Preparing changes in a separate address space
- Incremental migration of threads



## Multiverse

### Efficient Semi-Dynamic Variability

*EuroSys'19*

- Preparing variants at compile time
- Efficient binding by run-time code patching



## Wait-free code patching

### Run-time code patching without interruption

*OSDI'20*

- Preparing changes in a separate address space
- Incremental migration of threads



## Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```



## Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

**GCC**

## Code Segment

```
spin_irq_lock:  
    cmp    <smp>, 0  
    je     .else  
    cli  
    call   spin_acquire  
    ret  
.else:  
    cli  
    ret
```





## Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

+ Multiverse  
GCC

## Code Segment

```
spin_irq_lock.smp=1:  
cli  
call spin_acquire  
ret
```

```
spin_irq_lock.smp=0:  
cli  
ret
```

```
spin_irq_lock:  
cmp    <smp>, 0  
je     .else  
cli  
call   spin_acquire  
ret  
.else:  
cli  
ret
```



## Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

var

func

callsite

+ Multiverse  
GCC

Multiverse  
Deskriptoren

## Code Segment

```
spin_irq_lock.smp=1:  
    cli  
    call spin_acquire  
    ret
```

```
spin_irq_lock.smp=0:  
    cli  
    ret
```

```
spin_irq_lock:  
    cmp    <smp>, 0  
    je     .else  
    cli  
    call  spin_acquire  
    ret  
.else:  
    cli  
    ret
```



## Initial geladenes Code Segment

```
foo:
    ...
    call    multiverse_commit
    ...
    call    spin_irq_lock
    ...
    ret
```

```
spin_irq_lock.smp=1:
    cli
    call    spin_acquire
    ret
```

```
spin_irq_lock.smp=0:
    cli
    ret
```

```
spin_irq_lock:
    cmp     <smp>, 0
    je      .else
    cli
    call    spin_acquire
    ret
.else:
    cli
    ret
```

Multiverse  
Deskriptoren



## Initial geladenes Code Segment

```
foo:
    ...
    call multiverse_commit
    ...
    call spin_irq_lock
    ...
    ret
```

```
spin_irq_lock.smp=1:
    cli
    call spin_acquire
    ret
```

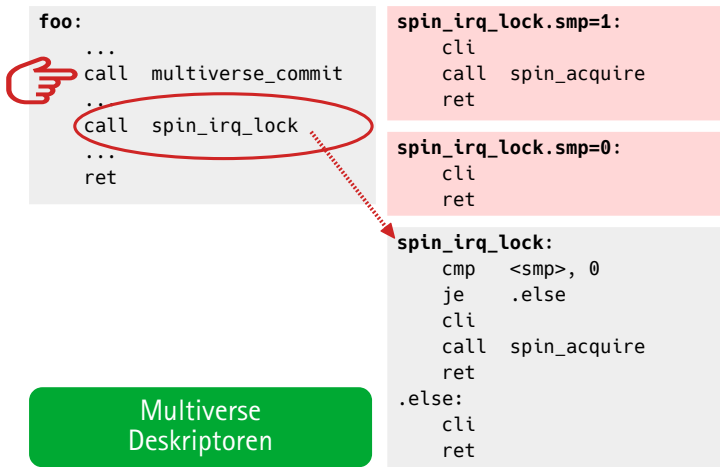
```
spin_irq_lock.smp=0:
    cli
    ret
```

```
spin_irq_lock:
    cmp    <smp>, 0
    je     .else
    cli
    call   spin_acquire
    ret
.else:
    cli
    ret
```

Multiverse  
Deskriptoren



## Initial geladenes Code Segment





Patched (smp == 1)

```
foo:
...
call multiverse_commit
...
call spin_irq_lock.smp=1
...
ret
```

```
spin_irq_lock.smp=1:
cli
call spin_acquire
ret
```

```
spin_irq_lock.smp=0:
cli
ret
```

```
spin_irq_lock:
jmp spin_irq_lock.smp=1
je .else
cli
call spin_acquire
ret
.else:
cli
ret
```

Multiverse  
Deskriptoren



Patched (smp == 0)

```
foo:
  ...
  call multiverse_commit
  ...
  cli
  ...
  ret
```

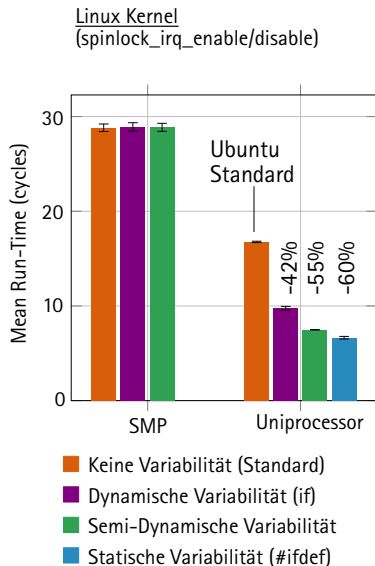
**Call-Site Inlining!**

```
spin_irq_lock.smp=1:
  cli
  call spin_acquire
  ret
```

```
spin_irq_lock.smp=0:
  cli
  ret
```

```
spin_irq_lock:
  jmp spin_irq_lock.smp=0
  je .else
  cli
  call spin_acquire
  ret
.else:
  cli
  ret
```

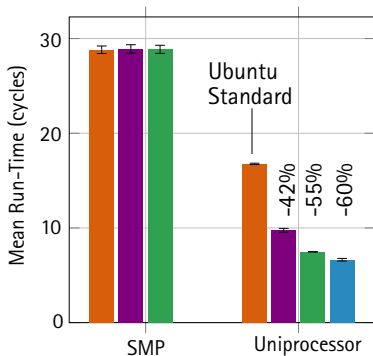
Multiverse  
Deskriptoren





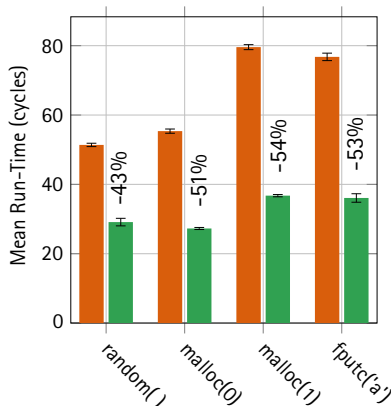


Linux Kernel  
(spinlock\_irq\_enable/disable)



- Keine Variabilität (Standard)
- Dynamische Variabilität (if)
- Semi-Dynamische Variabilität
- Statische Variabilität (#ifdef)

Musl C Bibliothek  
(Single Threaded Modus)

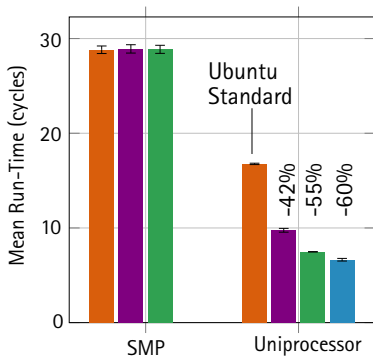


- Keine Variabilität (Standard)
- Semi-Dynamische Variabilität



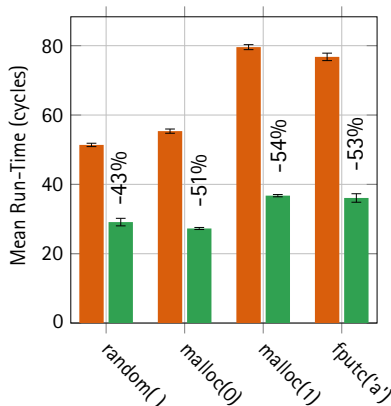
## Lock Elision in Kernel 4.16:

- 1161 spin-lock call sites
- +40 KiB size (zipped total: 10 MiB)



- Keine Variabilität (Standard)
- Dynamische Variabilität (if)
- Semi-Dynamische Variabilität
- Statische Variabilität (#ifdef)

## Musl C Bibliothek (Single Threaded Modus)



- Keine Variabilität (Standard)
- Semi-Dynamische Variabilität



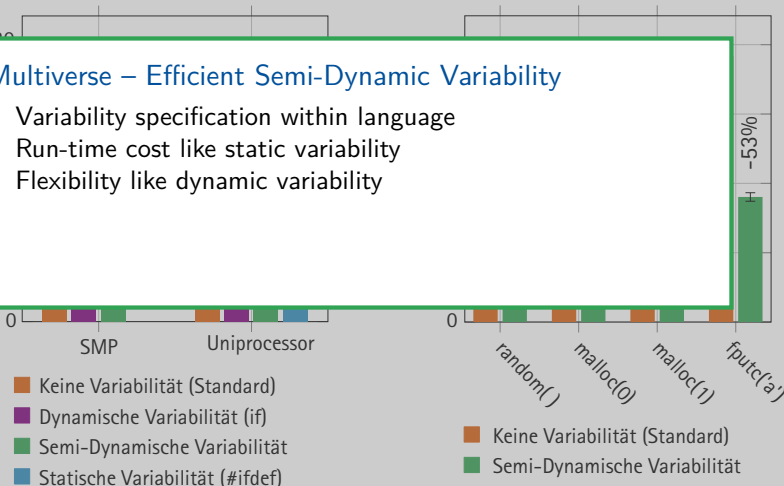
## Lock Elision in Kernel 4.16:

- 1161 spin-lock call sites
- +40 KiB size (zipped total: 10 MiB)

Musl C Bibliothek  
(Single Threaded Modus)

## Multiverse – Efficient Semi-Dynamic Variability

- + Variability specification within language
- + Run-time cost like static variability
- + Flexibility like dynamic variability





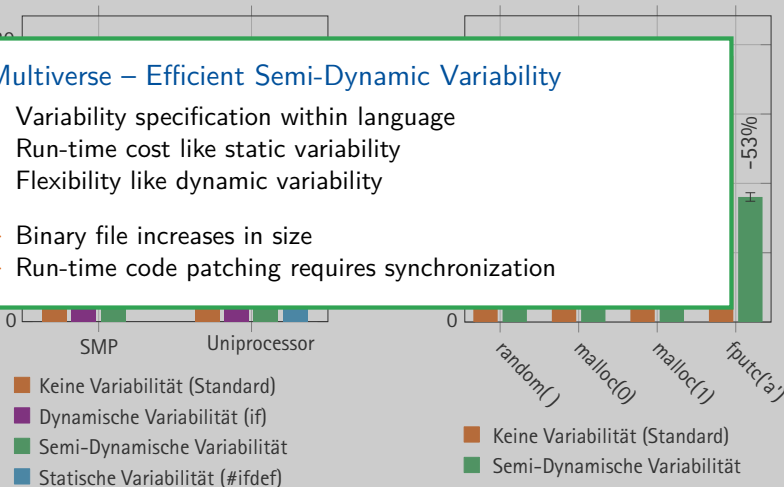
## Lock Elision in Kernel 4.16:

- 1161 spin-lock call sites
- +40 KiB size (zipped total: 10 MiB)

Musl C Bibliothek  
(Single Threaded Modus)

## Multiverse – Efficient Semi-Dynamic Variability

- + Variability specification within language
- + Run-time cost like static variability
- + Flexibility like dynamic variability
- Binary file increases in size
- Run-time code patching requires synchronization





## ■ Example Source Code

- `git clone https://github.com/luhsra/multiverse-atlas-demo.git`
- `./init =>` Clone dependent Libraries
- You need a mmview-enabled kernel for example 02/03

## ■ Virtual Machine Image

- Download: FIXME
- Start: `qemu-system-x86_64 -m 4G -enable-kvm -hda hda.qcow2 -smp 4`
- User/Passwort: root/root, user/user

## ■ Running Virtual Machine (temporary)

- `ssh user@lab.sra.uni-hannover.de -p 2250`
- Password: FIXME
- Machine will be available only for the time of the tutorial.



## Goal

### Demonstrate Semi-Dynamic Lock-Elision with Multiverse

- Static Creation with `__attribute__((multiverse))`
  - Multiverse Variables (MV) are „features“. Default Range: {0, 1}
  - Attributed functions are multiversed at MV usages.
  - Cross-Product of MV ranges.
- Binary File: 01-multiverse
  - Regular functions: `lock()`, `unlock()`
  - Use `objdump -d 01-multiverse` to inspect assembler
  - Search for `lock.multiverse.smp_0`

### ■ Multiverse Metadata

```
fn: lock 0x55bd78dc92e0, 2 variants
mvfn: 0x55bd78dc9410 (vars 1)
  assign: smp in [0, 0]
mvfn: 0x55bd78dc9420 (vars 1)
  assign: smp in [1, 1]
```



```
Without Multiverse (smp=0)
work() -> 83.376870 ms
work() -> 70.550493 ms
work() -> 70.670475 ms
work() -> 70.578593 ms
work() -> 73.273835 ms
Average (n=5): 73.690053
```

```
Without Multiverse (smp=1)
work() -> 511.101093 ms
work() -> 511.533894 ms
work() -> 512.103367 ms
work() -> 519.490820 ms
work() -> 511.268393 ms
Average (n=5): 513.099513
```

```
With Multiverse (smp=0)
work() -> 18.856244 ms
work() -> 18.820774 ms
work() -> 18.878794 ms
work() -> 18.757183 ms
work() -> 18.795324 ms
Average (n=5): 18.821664
```

```
With Multiverse (smp=1)
work() -> 509.799176 ms
work() -> 509.639876 ms
work() -> 518.221404 ms
work() -> 511.009342 ms
work() -> 510.144087 ms
Average (n=5): 511.762777
```



## Multiverse

Efficient Semi-Dynamic Variability

*EuroSys'19*

- Preparing variants at compile time
- Efficient binding by run-time code patching



## Wait-free code patching

Run-time code patching without interruption

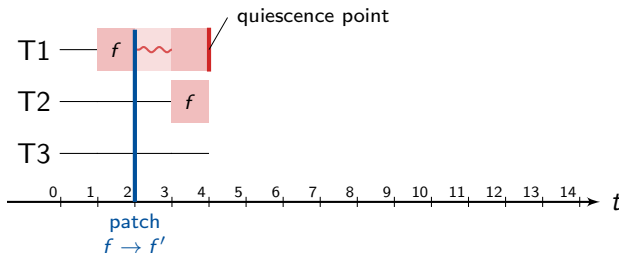
*OSDI'20*

- Preparing changes in a separate address space
- Incremental migration of threads



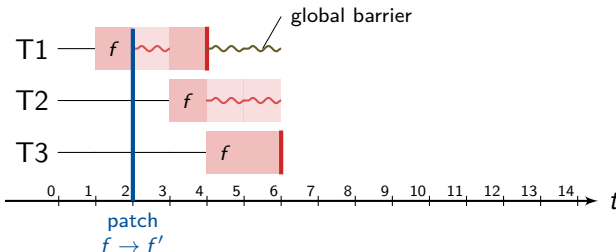


Run-Time Code Patching requires Quiescence!



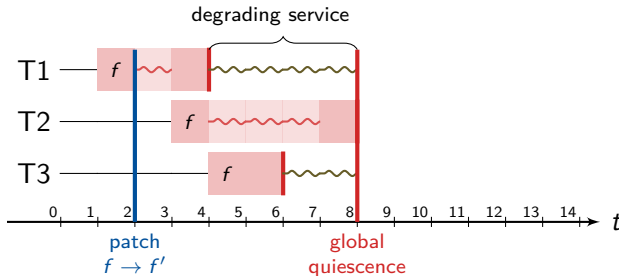


Run-Time Code Patching requires Quiescence!



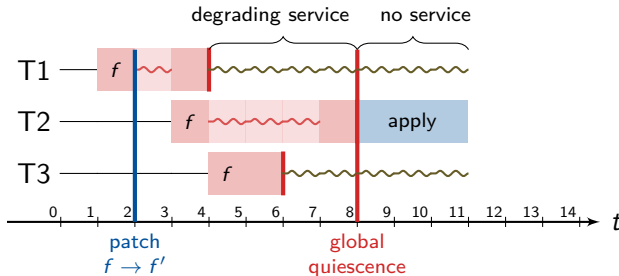


Run-Time Code Patching requires Quiescence!



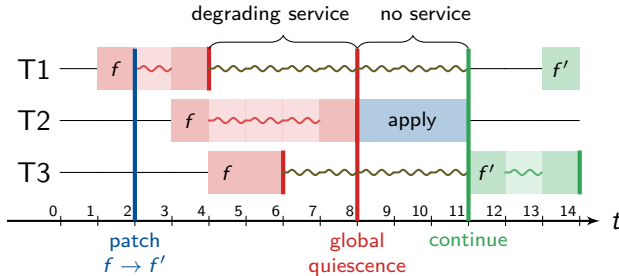
- Global barrier bears risk of dead-lock
- Impact on service quality

## Run-Time Code Patching requires Quiescence!



- Global barrier bears risk of dead-lock
- Impact on service quality
- + Strong consistency through global quiescence

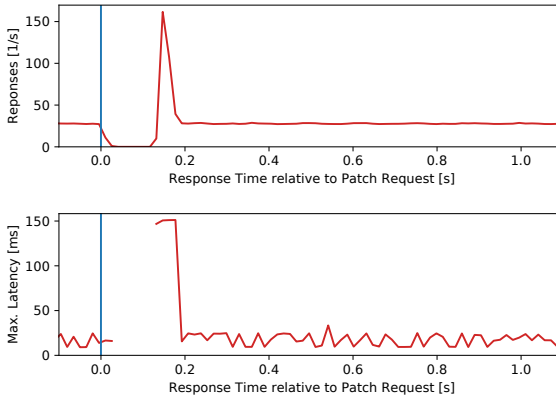
## Run-Time Code Patching requires Quiescence!



- Global barrier bears risk of dead-lock
- Impact on service quality
- + Strong consistency through global quiescence

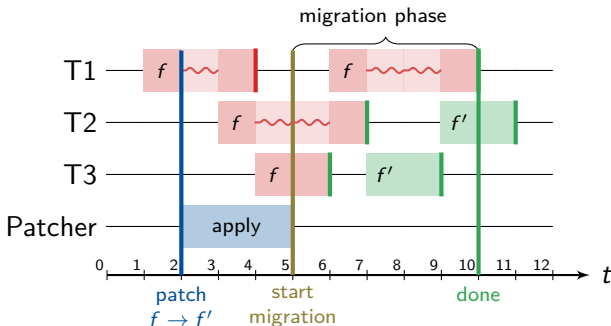


## OpenLDAP



- Glo
- Imp

+ Strong consistency through global quiescence



## Wait-Free Code Patching

## *Incremental Application of Patch*

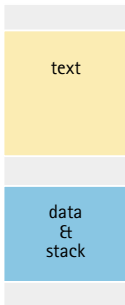
- Prepare the modified code segment
- Threads migrate themselves and without synchronization
- Old and new variants co-exist (for shorter or longer periods of time)



# Address-Space Views

address  
space

## Generation 0



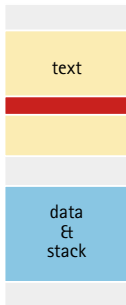




# Address-Space Views

address  
space

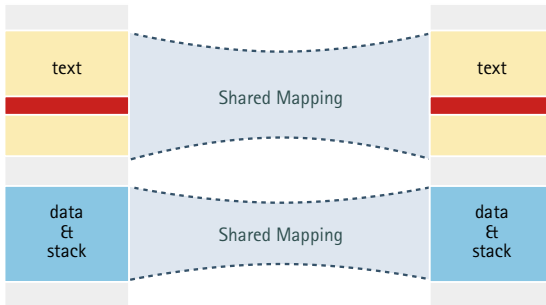
## Generation 0

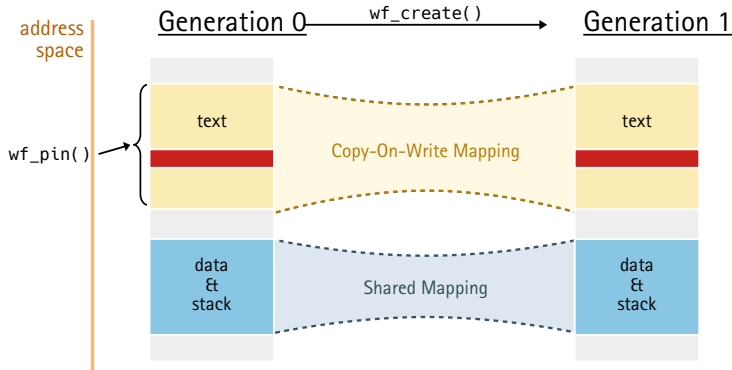


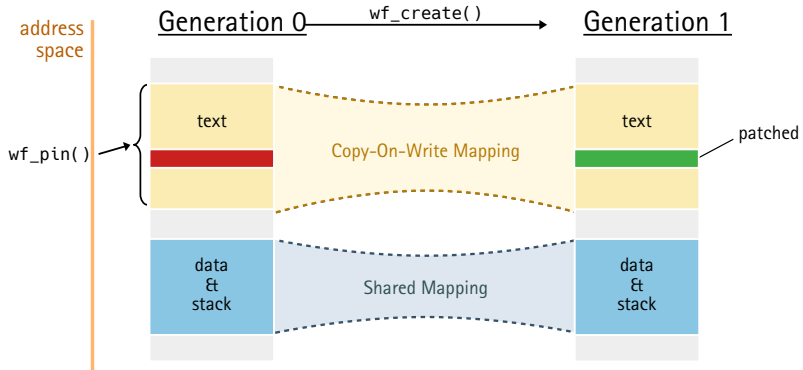


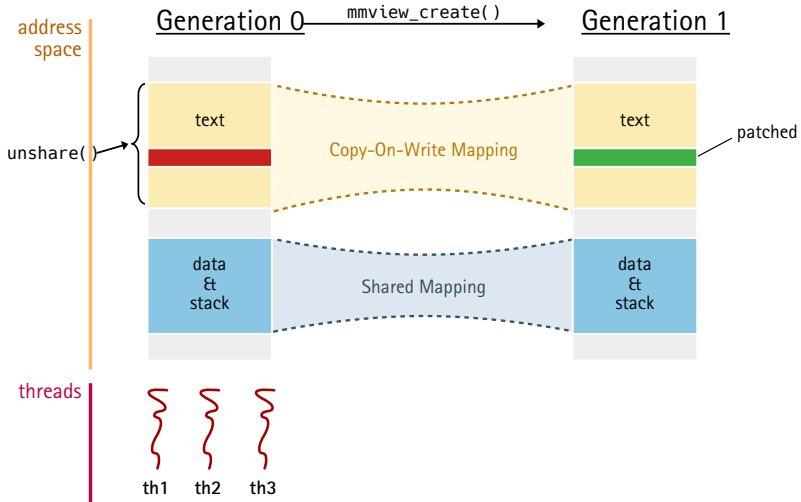
address  
space

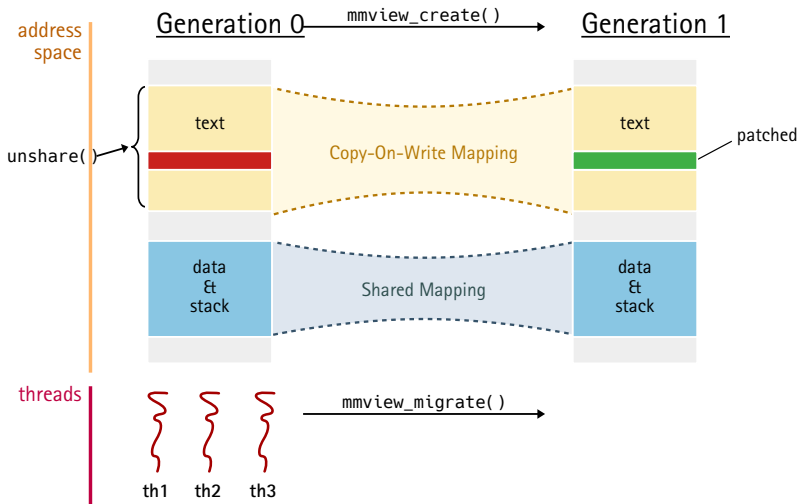
Generation 0  $\xrightarrow{\text{wf\_create()}}$  Generation 1

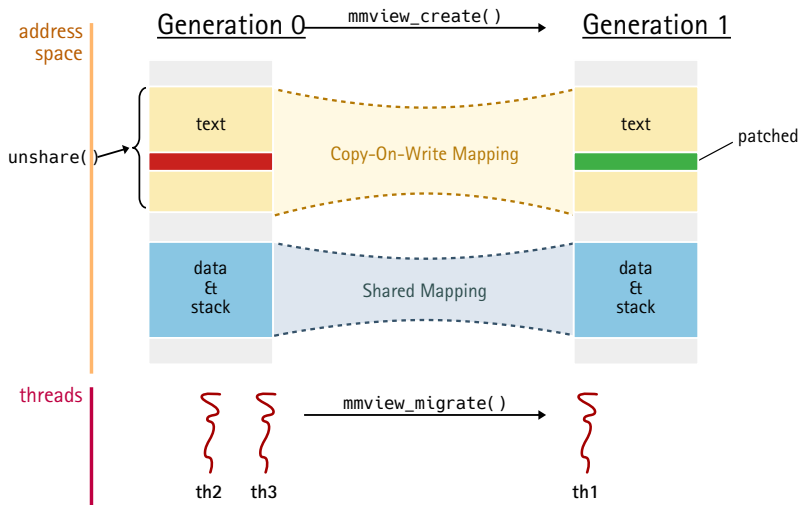


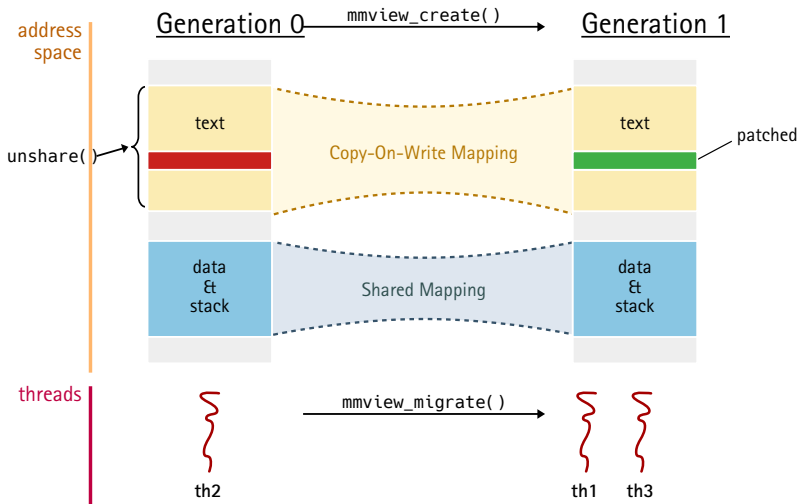




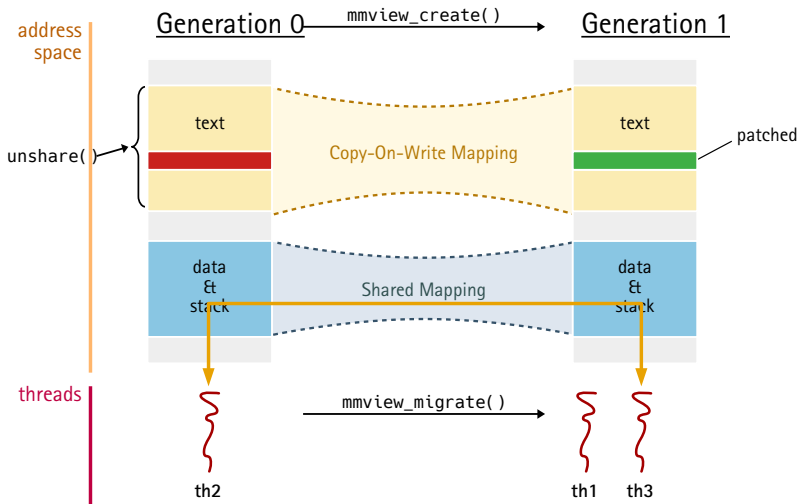






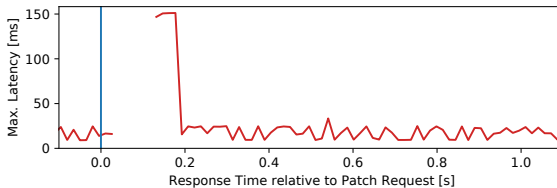
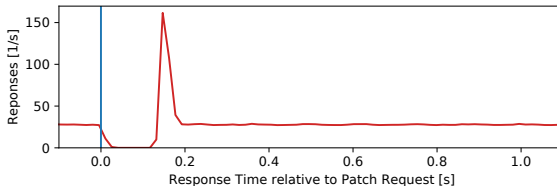








## OpenLDAP



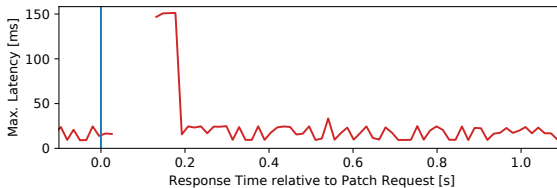
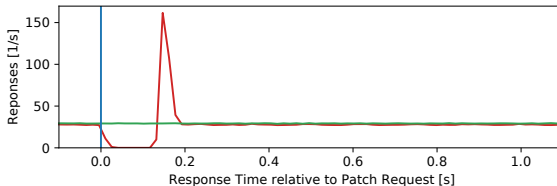
th2

th1

th3



## OpenLDAP



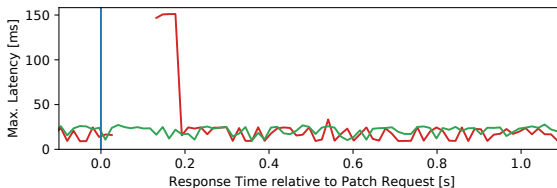
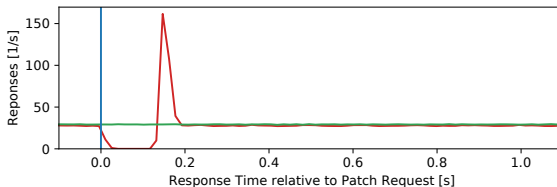
th2

th1

th3



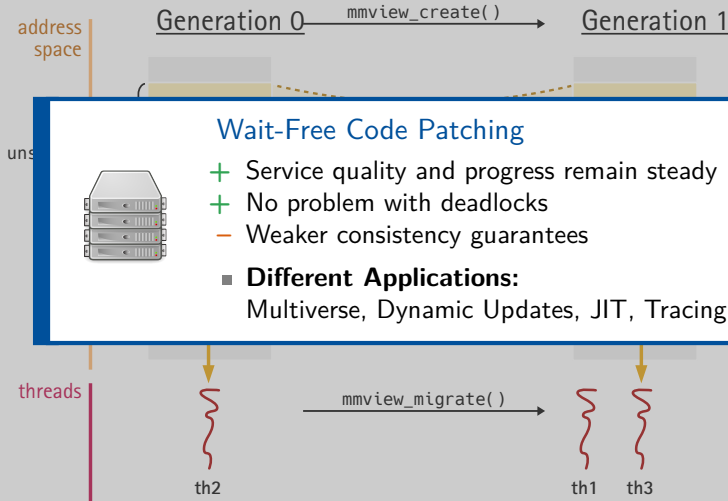
## OpenLDAP



th2

th1

th3





## Goal

### Demonstrate Address-Space View Migration with Multiverse

```
mmview_t view = mmview_create();  
mmview_migrate(view);  
  
multiversed_var = 1;  
multiverse_commit();  
desired_view = view;
```

- Create new address-space view and migrate the current thread.
- Commit current MV state into text segment.
- Warning: multiverse\_var is shared!

```
if (mmview_current() != desired_view) {  
    mmview_migrate(desired_mmview);  
}
```

Each thread checks at its local-quiescence point, whether it is already in the correct view. If not  
→ `mmview_migrate()`;



```
thread2 -> view B
thread0 -> view B
mmview_create() -> 1
mmview_migrate(1) -> 0
thread1 -> view B
thread1 migrate: mmview_migrate(1) -> 0
thread2 -> view B
thread2 migrate: mmview_migrate(1) -> 0
thread1 -> view A
thread0 -> view B
thread0 migrate: mmview_migrate(1) -> 0
thread1 -> view A
thread0 -> view A
thread2 -> view A
thread1 -> view A
thread0 -> view A
thread1 -> view A
thread1 -> view A
```



## Goal

Demonstrate concurrent existence of multiversed address-space views.

- **Problem:** Profiling in multi-threaded processes is slow!
  - One (shared!) counter per profiling point.
  - Massive cache line bouncing between threads.
  - Your Choice: Racy counter updates OR run-time impact of atomic operations
- **Approach:** Let only one thread profile at once
  - Two Views with Multiverse: Profiling and Non-Profiling View
  - Let N threads (probabilistically) switch to the profiling view
  - Control parallelism in profiling with semaphore





## ■ 4 threads calculating, no profiling

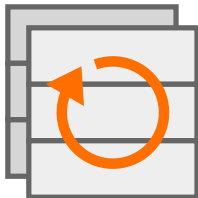
```
./03-concurrent 4 0  
Avg. time for fib(30) = 0.956386 ms; -nan branches/call 400
```

## ■ 4 threads calculating, 4 threads profiling

```
./03-concurrent 4 4  
Avg. time for fib(30) = 106.137128 ms; 0.509816 branches/call 400
```

## ■ 4 threads calculating, 1 threads profiling

```
./03-concurrent 4 1  
Avg. time for fib(30) = 1.088290 ms; 0.500000 branches/call 400
```



*Operating  
System  
Group*

- Operating System Group: Young Basic Research Group
  - Topics: Non-Volatile Memory, Variability, Dependable Real-Time Systems
  - Techniques: Virtual memory, Compilers, Fault Injection
  - Open PhD Position: 100% E13, 3 years
- Technical University Hamburg
  - ~8000 students, ~90 professors
  - Strong systems-oriented computer science
  - South of the Elbe

⇒ Get in contact, if interested



## Multiverse

Efficient Semi-Dynamic Variability

*EuroSys'19*

- Preparing variants at compile time
- Efficient binding by run-time code patching



## Wait-free code patching

Run-time code patching without interruption

*OSDI'20*

- Preparing changes in a separate address space
- Incremental migration of threads



- [1] Florian Rommel, Christian Dietrich, Daniel Friesel u. a. „From Global to Local Quiescence: Wait-Free Code Patching of Multi-Threaded Processes“. In: *14th Symposium on Operating System Design and Implementation (OSDI '20)*. Nov. 2020, S. 651–666. URL: <https://www.usenix.org/conference/osdi20/presentation/rommel>.
- [2] Florian Rommel, Christian Dietrich, Michael Rodin u. a. „Multiverse: Compiler-Assisted Management of Dynamic Variability in Low-Level System Software“. In: *Fourteenth EuroSys Conference 2019 (EuroSys '19)* (Dresden, Germany). New York, NY, USA: ACM Press, 2019. ISBN: 978-1-4503-6281-8. DOI: 10.1145/3302424.3303959.
- [3] Florian Rommel, Lennart Glauer, Christian Dietrich u. a. „Wait-Free Code Patching of Multi-Threaded Processes“. In: *Proceedings of the 10th SOSP Workshop on Programming Languages and Operating Systems (PLOS '19)*. Huntsville, ON, Canada: ACM, 2019, S. 23–29. ISBN: 978-1-4503-7017-2. DOI: 10.1145/3365137.3365404.
- [4] Valentin Rothberg, Christian Dietrich, Alexander Graf u. a. „Function Multiverses for Dynamic Variability“. In: *Foundations and Applications of Self\* Systems*. Hrsg. von Jesper Andersson, Rafael Capilla und Holger Eichelberger. Augsburg, 2016. URL: [https://www4.cs.fau.de/Publications/2016/rothberg\\_16\\_dspl.pdf](https://www4.cs.fau.de/Publications/2016/rothberg_16_dspl.pdf).