

Multiverse: Compiler-Assisted Management of Dynamic Variability in Low-Level System Software

Florian Rommel Christian Dietrich Michael Rodin Daniel Lohmann

{rommel, dietrich, lohmann}@sra.uni-hannover.de michael@rodin.online

Leibniz Universität Hannover, Germany

2019-03-28

Static Variability in Linux


```
void __init sched_init(void)
{
    int i, j;
    unsigned long alloc_size = 0, ptr;

    wait_bit_init();

#ifdef CONFIG_FAIR_GROUP_SCHED
    alloc_size += 2 * nr_cpu_ids * sizeof(void **);
#endif
#ifdef CONFIG_RT_GROUP_SCHED
    alloc_size += 2 * nr_cpu_ids * sizeof(void **);
#endif
    if (alloc_size) {
        ptr = (unsigned long)kzalloc(alloc_size, GFP_NOWAIT);

#ifdef CONFIG_FAIR_GROUP_SCHED
        root_task_group.se = (struct sched_entity **)ptr;
        ptr += nr_cpu_ids * sizeof(void **);

        root_task_group.cfs_rq = (struct cfs_rq **)ptr;
        ptr += nr_cpu_ids * sizeof(void **);
#endif
        /* CONFIG_FAIR_GROUP_SCHED */
#ifdef CONFIG_RT_GROUP_SCHED
        root_task_group.rt_se = (struct sched_rt_entity **)ptr;
        ptr += nr_cpu_ids * sizeof(void **);
#endif
    }
}
```



Static Variability in Linux

```
void __init sched_init(void)
{
    int i, j;
    unsigned long alloc_size = 0, ptr;


    wait_bit_init();

#ifdef CONFIG_FAIR_GROUP_SCHED
    alloc_size += 2 * nr_cpu_ids * sizeof(void **);
#endif
#ifdef CONFIG_RT_GROUP_SCHED
    alloc_size += 2 * nr_cpu_ids * sizeof(void **);
#endif
    if (alloc_size) {
        ptr = (unsigned long)kzalloc(alloc_size, GFP_NOWAIT);

#ifdef CONFIG_FAIR_GROUP_SCHED
        root_task_group.se = (struct sched_entity **)ptr;
        ptr += nr_cpu_ids * sizeof(void **);


        root_task_group.cfs_rq = (struct cfs_rq **)ptr;
        ptr += nr_cpu_ids * sizeof(void **);
#endif
        /* CONFIG_FAIR_GROUP_SCHED */
#ifdef CONFIG_RT_GROUP_SCHED
        root_task_group.rt_se = (struct sched_rt_entity **)ptr;
        ptr += nr_cpu_ids * sizeof(void **);
#endif
    }
}
```

Linux 5.0:
43918 #ifdefs





Dynamic Variability in Linux

Example: Operations for Paravirtualized Kernels (PV-Ops)





Dynamic Variability in Linux

Example: Operations for Paravirtualized Kernels (PV-Ops)




Dynamic Variability in Linux

Example: Operations for Paravirtualized Kernels (PV-Ops)




Dynamic Variability in Linux

Example: Operations for Paravirtualized Kernels (PV-Ops)




- Inside paravirtualization:
Privileged operations must be replaced by calls to the hypervisor
(e.g., enable/disable interrupts)




Dynamic Variability in Linux

Example: Operations for Paravirtualized Kernels (PV-Ops)




- Inside paravirtualization:
Privileged operations must be replaced by calls to the hypervisor
(e.g., enable/disable interrupts)
- Implemented by function pointers




Dynamic Variability in Linux

Example: Operations for Paravirtualized Kernels (PV-Ops)




- Inside paravirtualization:
Privileged operations must be replaced by calls to the hypervisor
(e.g., enable/disable interrupts)
- Implemented by function pointers → too much overhead




Dynamic Variability in Linux

Example: Operations for Paravirtualized Kernels (PV-Ops)




- Inside paravirtualization:
Privileged operations must be replaced by calls to the hypervisor
(e.g., enable/disable interrupts)
- Implemented by function pointers → too much overhead
- Run-time binary patching:
Replace indirect calls by direct calls



Dynamic Variability in Linux

arch/x86/include/asm/paravirt_types.h

```
349 #define _paravirt_alt(insn_string, type, clobber) \
350     "771:\n\t" insn_string "\n" "772:\n" \
351     ".pushsection .parainstructions,\"a\"\n" \
352     _ASM_ALIGN "\n" \
353     _ASM_PTR " 771b\n" \
354     "    .byte " type "\n" \
355     "    .byte 772b-771b\n" \
356     "    .short " clobber "\n" \
357     ".popsection\n"
...
360 #define paravirt_alt(insn_string) \
361     _paravirt_alt(insn_string, "%c[paravirt_tynenum]", \
    ↪         "%c[paravirt_clobber]")
...
570 asm volatile(pre \
571     paravirt_alt(PARAVIRT_CALL) \
572     post \
573     : call_clbr, ASM_CALL_CONSTRAINT \
574     : paravirt_type(op), \
575     paravirt_clobber(clbr), \
576     ##__VA_ARGS__ \
577     : "memory", "cc" extra_clbr);
```



Dynamic Variability in Linux


```
arch/x86/include/asm/paravirt_types.h
```

```
349 | #define paravirt_alt(insn_string_type clobber) \
```

- Complex implementation:

PV-Ops: 7 files, ~2000 loc (for x86)

```
572 | post  
573 | : call_clbr, ASM_CALL_CONSTRAINT  
574 | : paravirt_type(op),  
575 |   paravirt_clobber(clbr),  
576 |   ##__VA_ARGS__  
577 | : "memory", "cc" extra_clbr);
```




Dynamic Variability in Linux

```
arch/x86/include/asm/paravirt_types.h
```

```
349 | #define paravirt_alt(insn_string_type clobber) \
```

- Complex implementation:
PV-Ops: 7 files, ~2000 loc (for x86)
- Highly architecture-dependent: Multiple implementations

```
572 | post  
573 | : call_clbr, ASM_CALL_CONSTRAINT  
574 | : paravirt_type(op),  
575 |   paravirt_clobber(clbr),  
576 |   ##__VA_ARGS__  
577 | : "memory", "cc" extra_clbr);
```




Dynamic Variability in Linux

```
arch/x86/include/asm/paravirt_types.h
```

```
349 | #define paravirt_alt(insn_string type clobber) \
```

- Complex implementation:
PV-Ops: 7 files, ~2000 loc (for x86)
- Highly architecture-dependent: Multiple implementations
- Highly problem-specific: Multiple implementations
e.g., alternative instructions, SMP alternatives

```
572 | post  
573 | : call_clbr, ASM_CALL_CONSTRAINT  
574 | : paravirt_type(op),  
575 |   paravirt_clobber(clbr),  
576 |   ##__VA_ARGS__  
577 | : "memory", "cc" extra_clbr);
```



Dynamic Variability in Linux

```
arch/x86/include/asm/paravirt_types.h
```

```
349 | #define paravirt_alt(insn_string type clobber) \
```

- Complex implementation:
PV-Ops: 7 files, ~2000 loc (for x86)
 - Highly architecture-dependent: Multiple implementations
 - Highly problem-specific: Multiple implementations
e.g., alternative instructions, SMP alternatives
- Means for efficient dynamic variability are rarely used

```
572 | post  
573 | : call_clbr, ASM_CALL_CONSTRAINT  
574 | : paravirt_type(op),  
575 |   paravirt_clobber(clbr),  
576 |   ##__VA_ARGS__  
577 | : "memory", "cc" extra_clbr);
```



Efficient Dynamic Variability

Problems:

- Default approach: Performance costs (e.g., branches, function pointers)
- Binary patching: Code complexity → maintenance costs

Efficient Dynamic Variability

Problems:

- Default approach: Performance costs (e.g., branches, function pointers)
- Binary patching: Code complexity → maintenance costs



Multiverse

Compiler-Assisted Dynamic Variability via Binary Patching

Efficient Dynamic Variability

Problems:

- Default approach: Performance costs (e.g., branches, function pointers)
- Binary patching: Code complexity → maintenance costs



Multiverse

Compiler-Assisted Dynamic Variability via Binary Patching

Language extension to express efficient dynamic variability

- Express binary patching via standard control flow modification (if, ...)
- Generic mechanism for function-level run-time patching

Efficient Dynamic Variability

Problems:

- Default approach: Performance costs (e.g., branches, function pointers)
- Binary patching: Code complexity → maintenance costs



Multiverse

Compiler-Assisted Dynamic Variability via Binary Patching

Language extension to express efficient dynamic variability

- Express binary patching via standard control flow modification (if, ...)
 - Generic mechanism for function-level run-time patching
- *Compiler plugin + small run-time library*

Example: Linux Lock Elision

Linux Spinlock Implementation (simplified):

CONFIG_SMP set in the build system

```
void spin_irq_lock(raw_spinlock_t *lock) {  
    #ifdef CONFIG_SMP  
        irq_disable();  
        spin_acquire(&lock);  
    #else  
        irq_disable();  
    #endif  
}
```

Example: Linux Lock Elision

Linux Spinlock Implementation (simplified):

CONFIG_SMP set in the build system

```
void spin_irq_lock(raw_spinlock_t *lock) {  
    #ifdef CONFIG_SMP  
        irq_disable();  
        spin_acquire(&lock);  
    #else  
        irq_disable();  
    #endif  
}
```

Example: Linux Lock Elision

Linux Spinlock Implementation (simplified):

CONFIG_SMP set in the build system

```
void spin_irq_lock(raw_spinlock_t *lock) {  
    #ifdef CONFIG_SMP  
        irq_disable();  
        spin_acquire(&lock);  
    #else  
        irq_disable();  
    #endif  
}
```

Example: Linux Lock Elision

```
bool smp;
```

```
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

Example: Linux Lock Elision

```
bool smp;
```

```
void spin_irq_lock(...) {
```

```
    if (smp) {
```

```
        irq_disable();
```

```
        spin_acquire(&lock);
```

```
    } else {
```

```
        irq_disable();
```

```
    }
```

```
}
```

branched control flow
run-time overhead



Example: Linux Lock Elision

```
bool smp;
```

```
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo(void) {  
    smp = true;
```

Example: Linux Lock Elision

```
__attribute__((multiverse))
```

```
bool smp;
```

```
__attribute__((multiverse))
```

```
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo(void) {  
    smp = true;
```

Example: Linux Lock Elision

```
__attribute__((multiverse))
```

```
bool smp;
```

```
__attribute__((multiverse))
```

```
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo(void) {  
    smp = true;  
    multiverse_commit();  
    // ...  
  
    spin_irq_lock(lock);  
}
```

Example: Linux Lock Elision

```
__attribute__((multiverse))
```

```
bool smp;
```

```
__attribute__((multiverse))
```

```
void spin_irq_lock(...) {
```

```
    if (smp) {
```

```
        irq_disable();  
        spin_acquire(&lock);
```

```
    } else {
```

```
        irq_disable();
```

```
    }
```

```
}
```

```
void foo(void) {
```

```
    smp = true;
```

```
    multiverse_commit();
```

```
    // ...
```

```
    spin_irq_lock(lock);
```

```
}
```

Example: Linux Lock Elision

```
__attribute__((multiverse))
```

```
bool smp;
```

```
__attribute__((multiverse))
```

```
void spin_irq_lock(...) {
```

```
    if (smp) {
```

```
        irq_disable();
```

```
        spin_acquire(&lock);
```

```
    } else {
```

```
        irq_disable();
```

```
    }
```

```
}
```

```
void foo(void) {
```

```
    smp = false;
```

```
    multiverse_commit();
```

```
    // ...
```

```
    spin_irq_lock(lock);
```

```
}
```

Ahead-of-Time Variant Generation

Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

Ahead-of-Time Variant Generation

Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

GCC

Code Segment

```
spin_irq_lock:  
    cmp    <smp>, 0  
    je     .else  
    cli  
    call   spin_acquire  
    ret  
.else:  
    cli  
    ret
```

Ahead-of-Time Variant Generation

Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

+ Multiverse
GCC

Code Segment

```
spin_irq_lock.smp=1:  
    cli  
    call spin_acquire  
    ret
```

```
spin_irq_lock.smp=0:  
    cli  
    ret
```

```
spin_irq_lock:  
    cmp    <smp>, 0  
    je     .else  
    cli  
    call  spin_acquire  
    ret  
.else:  
    cli  
    ret
```


Ahead-of-Time Variant Generation

Quellcode

```
__attribute__((multiverse))  
bool smp;
```

```
__attribute__((multiverse))  
void spin_irq_lock(...) {  
    if (smp) {  
        irq_disable();  
        spin_acquire(&lock);  
    } else {  
        irq_disable();  
    }  
}
```

```
void foo() {  
    //...  
    spin_irq_lock();  
    //...  
}
```

var

func

callsite

Multiverse
+
GCC

Multiverse
Deskriptoren

Code Segment

```
spin_irq_lock.smp=1:  
    cli  
    call spin_acquire  
    ret
```

```
spin_irq_lock.smp=0:  
    cli  
    ret
```

```
spin_irq_lock:  
    cmp    <smp>, 0  
    je     .else  
    cli  
    call  spin_acquire  
    ret  
.else:  
    cli  
    ret
```

Run-Time Patching

Initial geladenes Code Segment

```
foo:
    ...
    call    multiverse_commit
    ...
    call    spin_irq_lock
    ...
    ret
```

```
spin_irq_lock.smp=1:
    cli
    call    spin_acquire
    ret
```

```
spin_irq_lock.smp=0:
    cli
    ret
```

```
spin_irq_lock:
    cmp     <smp>, 0
    je      .else
    cli
    call    spin_acquire
    ret
.else:
    cli
    ret
```

Multiverse
Deskriptoren

Run-Time Patching

Initial geladenes Code Segment

foo:

```
...  
call multiverse_commit  
..  
call spin_irq_lock  
...  
ret
```

spin_irq_lock.smp=1:

```
cli  
call spin_acquire  
ret
```

spin_irq_lock.smp=0:

```
cli  
ret
```

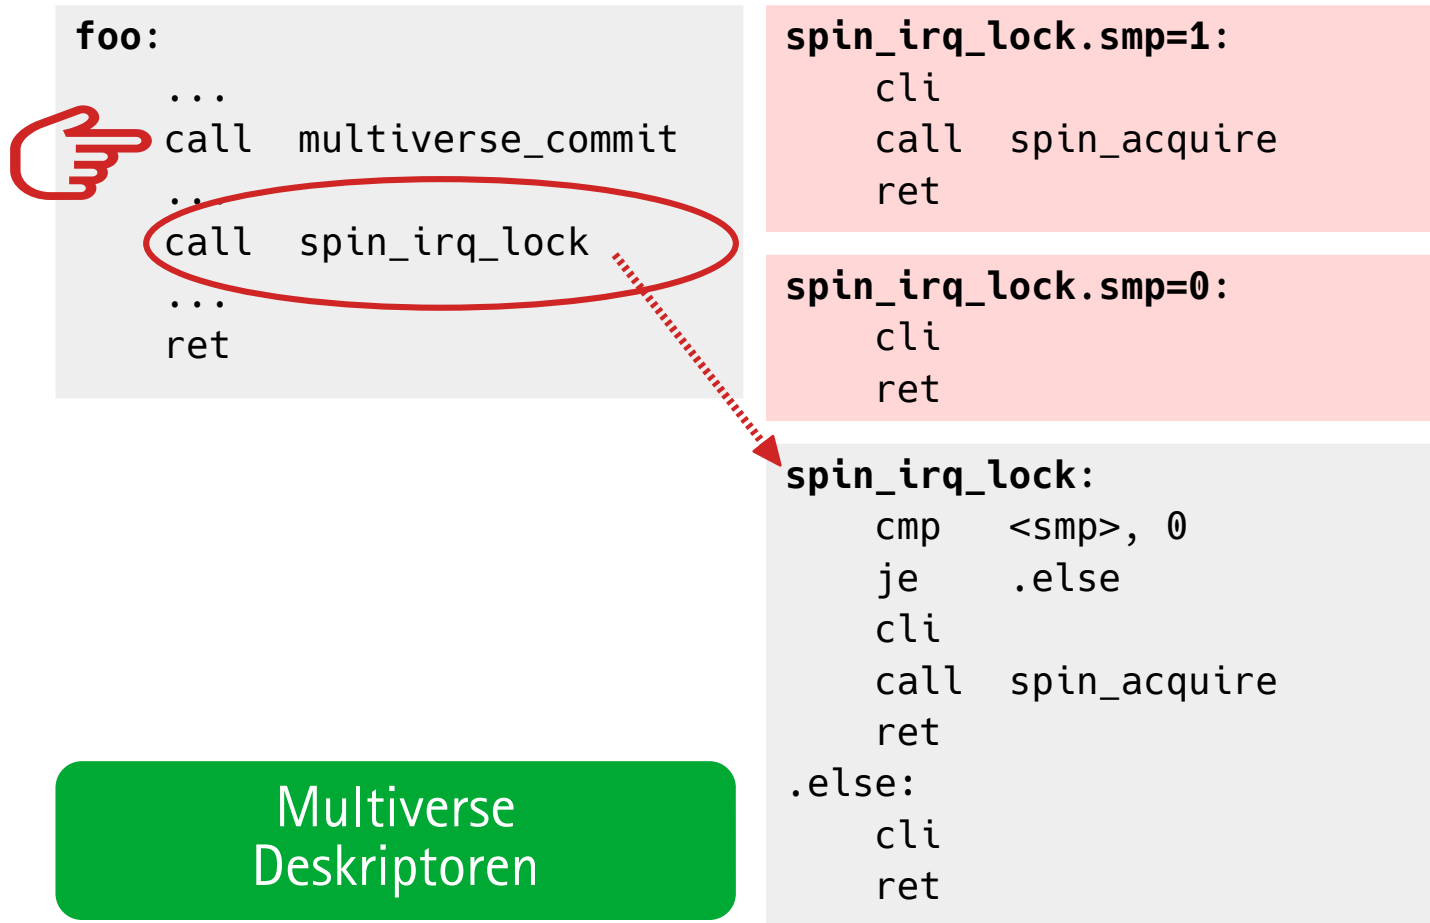
spin_irq_lock:

```
cmp    <smp>, 0  
je     .else  
cli  
call   spin_acquire  
ret  
.else:  
cli  
ret
```

Multiverse
Deskriptoren

Run-Time Patching

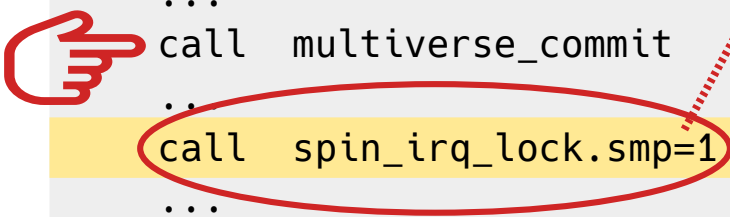
Initial geladenes Code Segment



Run-Time Patching

Patched (smp == 1)

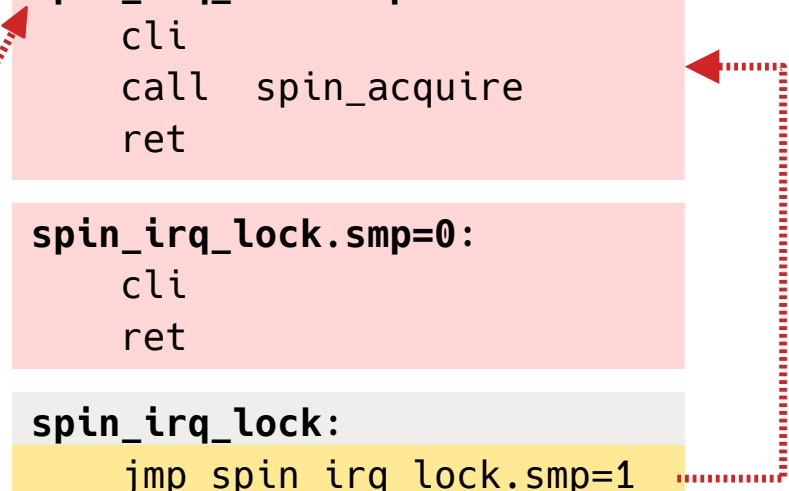
```
foo:
...
call multiverse_commit
..
call spin_irq_lock.smp=1
...
ret
```



```
spin_irq_lock.smp=1:
cli
call spin_acquire
ret
```

```
spin_irq_lock.smp=0:
cli
ret
```

```
spin_irq_lock:
jmp spin_irq_lock.smp=1
je .else
cli
call spin_acquire
ret
.else:
cli
ret
```




Multiverse
Deskriptoren

Run-Time Patching

Patched (smp == 1)

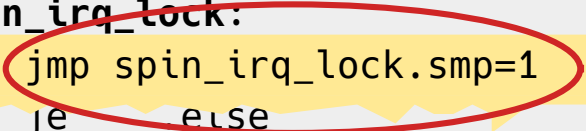
```
foo:
...
call multiverse_commit
...
call spin_irq_lock.smp=1
...
ret
```



```
spin_irq_lock.smp=1:
cli
call spin_acquire
ret
```

```
spin_irq_lock.smp=0:
cli
ret
```

```
spin_irq_lock:
jmp spin_irq_lock.smp=1
je .else
cli
call spin_acquire
ret
.else:
cli
ret
```




Multiverse
Deskriptoren

Run-Time Patching

Patched (smp == 1)

```
foo:
...
call multiverse_commit
...
call spin_irq_lock.smp=1
...
ret
```



```
spin_irq_lock.smp=1:
cli
call spin_acquire
ret
```

```
spin_irq_lock.smp=0:
cli
ret
```

```
spin_irq_lock:
jmp spin_irq_lock.smp=1
je .else
cli
call spin_acquire
ret
.else:
cli
ret
```

Multiverse
Deskriptoren

Run-Time Patching

Patched (smp == 0)

```
foo:
  ...
  call multiverse_commit
  ...
  cli
  ...
  ret
```

Call-Site Inlining!

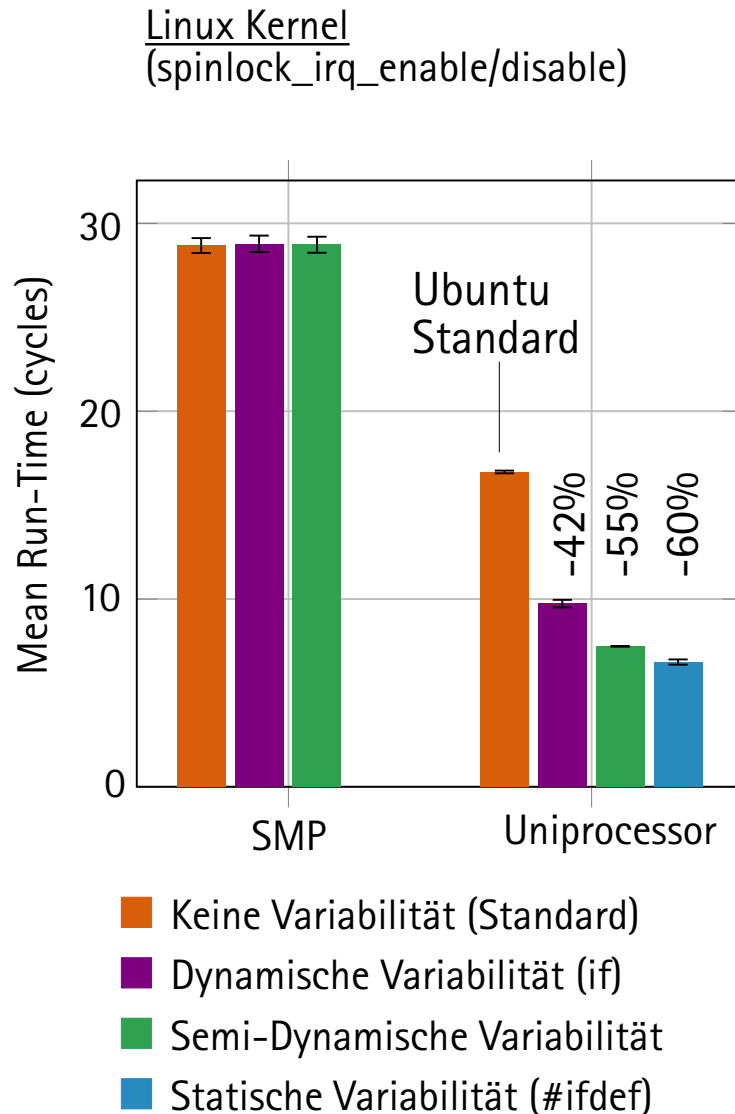
```
spin_irq_lock.smp=1:
  cli
  call spin_acquire
  ret
```

```
spin_irq_lock.smp=0:
  cli
  ret
```

```
spin_irq_lock:
  jmp spin_irq_lock.smp=0
  je .else
  cli
  call spin_acquire
  ret
.else:
  cli
  ret
```


Multiverse
Deskriptoren

Evaluation: Lock Elision in Kernel und Userspace




Evaluation: Lock Elision in Kernel und Userspace

Linux Kernel
(spinlock_irq_enable/disable)



- Keine Variabilität (Standard)
- Dynamische Variabilität (if)
- Semi-Dynamische Variabilität
- Statische Variabilität (#ifdef)

Musl C Bibliothek
(Single Threaded Modus)




- Keine Variabilität (Standard)
- Semi-Dynamische Variabilität

Evaluation: Lock Elision in Kernel und Userspace


Lock Elision in Kernel 4.16:

- 1161 spin-lock call sites
- +40 KiB size (zipped total: 10 MiB)




- Keine Variabilität (Standard)
- Dynamische Variabilität (if)
- Semi-Dynamische Variabilität
- Statische Variabilität (#ifdef)

Musl C Bibliothek (Single Threaded Modus)




- Keine Variabilität (Standard)
- Semi-Dynamische Variabilität

Evaluation – Userspace Microbenchmarks




Evaluation – Userspace Microbenchmarks

musl libc – Lock Elision
(Ein Thread)



(Mehrere Threads)



■ no multiverse (default) ■ multiverse

Evaluation

■ GNU Grep

→ Optimized for more than 30 years

*Multiversed a conditional branch in the inner loop
(recognition of multi-byte characters on/off)*

Evaluation

■ GNU Grep

→ Optimized for more than 30 years

*Multiversed a conditional branch in the inner loop
(recognition of multi-byte characters on/off)*

→ End-to-end measurement: **-2.73 % run-time**

→ Only 50 changed lines to use multiverse

Evaluation

■ GNU Grep

→ Optimized for more than 30 years

*Multiversed a conditional branch in the inner loop
(recognition of multi-byte characters on/off)*

→ End-to-end measurement: **-2.73 % run-time**

→ Only 50 changed lines to use multiverse

■ Multiverse – Numbers

GCC Plugin

Lines of Code: <1200
+ compatibility headers:
~1300

GCC version: 6.3 and higher

Run-Time Library

Lines of Code: < 850
Compiled: 6.5 KiB
Architectures: IA-32, AMD64
ARM [to come]

Summary



Multiverse

Compiler-Assisted Dynamic Variability via Binary Patching

- Language extension for easy-to-use, efficient dynamic variability

GCC-Plugin: Generates specialized function variants

Run-Time Library: Function-level binary patching

- Evaluation
 - Consolidation of current patching mechanisms (PV-Ops)
 - Introduction of new dynamic variation points (Lock Elision, Grep)

- Try it: <https://github.com/luhsra/multiverse>

